TLM

SCHOOL OF COMPUTATION, INFORMATION AND TECHNOLOGY — INFORMATICS

TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

GPU Coroutines in Autonomous Driving

Jaden Rotter

TUM

SCHOOL OF COMPUTATION, INFORMATION AND TECHNOLOGY — INFORMATICS

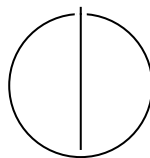TECHNISCHE UNIVERSITÄT MÜNCHEN

Bachelor's Thesis in Informatics

GPU Coroutines in Autonomous Driving

GPU Coroutines in Autonomes Fahren

Author:	Jaden Rotter
Examiner:	Supervisor
Supervisor:	Jianfeng Gu
Submission Date:	Submission date

I confirm that this bachelor's thesis is my own work and I have documented all sources and material used.

Munich, Submission date   Jaden Rotter

Acknowledgments

Abstract

Contents

1 Introduction

Autonomous driving presents a formidable challenge in computing, requiring complex high-performance workloads that heavily rely on **GPU!**s, yet traditional scheduling techniques fail to deliver the necessary predictability and reliability demanded by real-time systems. **GPU! (GPU!)** hardware is optimized for maximizing throughput, a focus that compromises low-latency performance critical for real-time applications. Similarly, traditional **GPU!** scheduling methods do not consider the strict deadlines essential for real time systems, particularily under conditions of **GPU!** contention, where multiple tasks compete for limited computational resources. This thesis will implemented an approach to **GPU!** scheduling using coroutines on persistant threads to reduce latency variability and address contention issues, ultimately delivering the predictability and reliability necessary for real-time autonomous driving applications.

1.0.1 Necessesity of GPU!s in Autonomous Driving

Recently, **GPU!**s have revolutionized data processing with their ability to handle massive amounts of data and execute highly-parallelized computations. Previous data processing systems, which used traditional CPUs, by contrast, relied on a **MIMD! (MIMD!)** architecture that excels at handling complex control logic at high frequencies by executing different instructions on seperate data streams concurrently. However, for tasks that require the repeated execution of singular instructions such as large-scale data processing workloads, CPUs are burdened by the overhead of the complex control logic, where a simpler, more parallel design would be more effective. **GPU!**s, created to overcome this limitation, use a **SIMT! (SIMT!)** architecture with thousands of simpler, slower threads executing simultaneously. The parallelism provided by the **SIMT!** model allows a far higher throughput that outperforms **CPU!**s. Here, each data element is processed by its own thread, allowing the same instruction to be executed simultaneously across numerous data streams.

Higher data processing performance is vital to autonomous systems, which rely extensively on machine learning tasks and sensor data processing. In autonomous systems a multitude of modules require deep learning, a machine learning technique, to automatically extract patterns, such as computer vision, and make decisions [**JEON2021167**]. Deep learning, inspired by the structure of the human brain, is composed of layers

of numerous interconnected, identical nodes called neural networks. Neural network models rely on mathematical operations between different neighboring layers to perform inference, essentially the prediction making or recognition process. The layers are represented as matrices, which then get multiplied and convoluted with one another and are used by specialized functions, called activation functions, to introduce non-linearity. With very large neural networks with thousands or millions of nodes, the inference computation is repetitive and slow. The capability to execute these repeptitive workloads concurrently across different dataset makes **GPU!**s fundamental to performance in tasks using complex neural networks. Furthermore, **GPU!**s are necessary to process the data rate produced by high-bandwidth, high-frequency sensors used in autonomous driving. The parallelism in **GPU!**s makes them the ideal platform over CPUs for deep learning and data processing tasks, tasks fundamental for autonomous driving systems.

1.0.2 Autonomous Driving as a Real Time System

Real time systems, such as autonomous driving, are designed with strict timing constraints in mind, to ensure predictable and deterministic behavior [**10155700**]. They use a concept of deadlines, which are subdivided into soft and hard-deadlines. Hard deadlines are critical and a failure leads to the systems failure or unsafe conditions. For example, in autonomous driving, collision avoidance with another vehicle and brake activation are hard deadlines. If these deadlines are not met, the safety of the passengers and the system is at risk. On the other hand, soft deadlines are not critical and missing these deadlines degrades performance, but does not cause system failure. In autonomous driving, this would show in route planning and navigation updates, where a delay would lead to suboptimal paths, but safety is not comprimised. Real time systems need to be capable of effectively and efficiently switching from lower priority tasks, soft deadline tasks, to high priority, hard deadline tasks, to ensure the safety both for the passengers and nearby individuals.

1.0.3 GPU weaknesses in Real Time Systems

Given the reliance on **GPU!**s in autonomous driving systems for data processing and machine learning, they need to meet the strict timing requirements of real-time applications, which is lacking under current scheduling methods. GPU tasks are traditionaly queued and scheduled based on availabilty to optimize for high throughput applications like graphics rendering and offline machine learning training. These tasks run to completion, similar to a batch system, meaning multiple tasks, such as the different modules in autonomous driving execute sequentially. Thus, depending on

the current workload for the GPU, the execution time and latency can vary, which poses a safety risk. To adapt **GPU!**s to the requirements of real time systems, while maintaining the advantages they provide in processing power, the scheduling of these systems needs to support asynchronous programming.

1.1 Background

1.1.1 Asynchronous Programming and Coroutines

Asynchronous programming is a method of programming a system to handle tasks concurrently instead of sequentially. Typically used in conjunction with tasks that delay or have high wait-times, such as I/O heavy jobs, asynchronous programming reduces overall execution time by more efficiently using processing ressources. For example, while waiting for I/O heavy input like sensor I/O, asynchronous code lets other tasks execute in the meantime, before returning when the data arrives. For real-time systems, asynchronous programming additionally uses the intermittant execution model to enforce determinism. By allowing the **GPU!**s to switch between concurrent tasks, hard deadlines can be immediately enforced without delay.

Coroutines, an implementation of asynchronous programming, uses suspendable functions to halt execution. Suspendable functions are implemented by capturing the current context, know as the continuation, of the currently running thread and save the data to be run later [**Zheng2022LuisaRender**]. After being saved, a new process can take over execution, without interrupting or overwritting the state of the previous process. Once the intermittant process or higher priority process has finished execution, the original task can continue executing by restoring the process context, which was previously saved. Capturing the continuation of a function allows the resumption of the program to be strategically deferred.

1.1.2 Autonomous Vehicles

Autonomous driving systems, such as self-driving cars, are subdivided into several modules, perception, localization, mapping, planning and control, which all rely on massive data processing and machine learning, to help the system to understand its enviornment, plan its trajectory and drive without any human help. Perception gathers and interprets vast amounts of data from numerous sensors, such as **LiDAR!** (**LiDAR!**) sensors, cameras, radar, and ultrasonic sensors, to identify and classify surrounding objects [**apollo**]. After perception, localization and mapping use a **SLAM!** (**SLAM!**) algorithm to build maps and determine the vehicle's precise location within its enviornment and the generated map. Using the foundation set by localization and

mapping, path planning determines the vehicles trajectory from its current location to a target destination, while respecting traffic laws and avoiding obstacles. Lastly, control actively follows the trajectory and sets the speed, steering and braking required for the planned path. Autonomous driving sytems are constantly computing each of these modules based on updates in the surroundings, which requires significant computational ressources to maintain a low latency.

Specifically, these modules require large amounts of data to accurately compute , which are based in deep learning, a machine learning technique that demands vast amounts of computational resources. In perception, deep learning models, particularily **CNN!**s identify objects, lane markings, pedestrians, vehicles, and traffic lights from the raw sensor data. Beyond perception, deep learning enhances localization by refining sensor fusion techniques, improving accuracy in determining the vehicles position. Furthermore, deep learning also supports planning and control, through reinforcement learning and predictive models to optimize trajectories. The real-time nature of these components are needed to reduce the latency, which requires these models to continously be updating their predictions as new sensor input arrives and the vehicle moves. This places a high demand on the hardware, requiring high-performance accelerators to process vast amounts of data efficiently. To meet these requirements, these computations are scheduled on to a **GPU!**, which provides the high degree of parallelism required by the vast amount of data and computations within autonomous driving.

1.1.3 NVIDIA Tesla V100 Architecture

For the purpose of this thesis, the NVIDIA Tesla V100 accelerator, which uses the Volta GV100 **GPU!** architecture, was selected due to its availability and high-performance computing capabilities. At the core of its execution model is the warp, a group of 32 threads executing in **SIMT!** fashion [**nvidia2017tesla**]. Each warp is scheduled and dispatched within one of the 4 processing partitions of a **SM!** (**SM!**). The warp scheduler and warp dispatcher each handle 32 threads per clock cycle. Within each **SM!** partition, there are Tensor Cores for deep learning, 64 bit-**FP!** (**FP!**) cores, **LD/ST!** (**LD/ST!**) units, a register file, and **SFU!**s for mathematical functions such as sine and square root. Additionally, each partition contains 16 Cuda cores, which can execute both **FP!** 32 bit and **INT!** (**INT!**) 32 bit instructions, but not simultaneously. Each partition has its own L0 instruction cache, while all partitions share a L1 instruction and data cache. A **SM!** can support up to 64 concurrent warps, allowing a maximum of $64 * 32 = 2048$ threads per **SM!**. At a higher level, two **SM!**s are grouped to form a **TPC!** (**TPC!**). The Tesla V100 GPU consists of six **GPC!**s, each containing seven **TPC!**s, resulting in a total of 84 **SM!**s across the chip. This hierarchical organization, from warps to **SM!**s, **TPC!**s,

and **GPC!**s, enables the Tesla V100 to efficiently handle massively parallel workloads, making it well-suited for high-performance computing and deel learning applications.

2 Objectives

The objective for this thesis is to develop, implement and evaluate a persistant thread with coroutine based scheduling approach for GPU threads in Apollo[1], an open source autonomous driving platform, to improve real-time safety and determinism by ensuring a predictable scheduling behavior. This requires analyzing the limitations of existing GPU scheduling techniques, which, due to their batch processing design, lack the ability to perform task switching. To adress this, the new coroutine based scheduling mechanism LuisaCompute-coroutine[2] will be employed. The Luisa coroutine scheduling was designed for graphics rendering tasks, around which their **DSL!** (**DSL!**) is written. Therefore, the scheduler and its **DSL!** will need to be adapted to fit the requirements of the autonomous driving domain, which prioritizes responsiveness, fault tolerance, and strict timing guarantees. By integrating this feature into Apollo, the GPU will be able to use coroutines for more flexible and efficient task management, ultimately enhancing the reliability and safety of the real-time autonomous driving platform.

2.0.1 Measurement and Evalutation

To measure the success of a coroutine based solution, the latency of new time sensitive task scheduling and deadline success will be evaluated in comparison to the old system. Furthermore, the implementability and effectiveness of this application will be analyzed through practical experiments and benchmarks, focusing on how well the system maintains deterministic behavior under varying workloads and task complexity.

---

[1]Apollo is an open-source project developed by Baidu. For more information, visit the GitHub repository: `https://github.com/ApolloAuto/apollo`

[2][**Zheng2022LuisaRender**] and the accompanying source code `https://github.com/LuisaGroup/LuisaCompute-coroutine/tree/next`

3 Literature Review

3.1 GPU Coroutines for Flexible Splitting and Scheduling of Rendering Tasks

Previous work, on which this thesis is based, offers a new method in rendering task management to demonstrate a flexible, fine-grained scheduling mechanism on **GPU!**s [**Zheng2022LuisaRender**]. Previously discussed in the background and introduction, this paper focuses on splitting tasks into discrete segments that can be suspended and resumed as needed on **GPU!**s. Specifically, large, monolithic kernels could be rewritten into smaller tasks using coroutines. This capability, important for tasks with downtimes, reduced the latency associated with waiting for an entire batch of rendering computation to complete and maximized the utilization of the **GPU!**'s parallel processing capabilities.

In this paper, a flexible **DSL!** was presented, which allowed the writing in a mega-kernel fashion with $suspend statements to define suspention marks, to define a coroutine. These points, allow the creation of different discrete segments that can be suspended and resumed as needed. The **DSL!** has support for a multitude of languages, specifically C# and CUDA. With minimal changes to the original code, the new DSL can be included, which becomes dynamicaly recorded and translated to an intermediate representation. From the intermediate representation, after a set of comiler transformation and analysis used to extract the state frame, a scheduler can define the execution and manage the state frames. Important here, is that the scheduler can be chosen or rewritten, which will be used for Apollo. Ultimately, through the design of an easy-to-use and easily accessible library, this work can be implemented into autonomous driving systems.

4 Proposed Approach

The proposed solution leverages LuisaCompute-coroutines to introduce a coroutine-based scheduling mechanism into Apollo. The approach begins by integrating the LuisaCompute-coroutine framework to manage GPU threads more flexibly, allowing rendering and real-time computation tasks to be suspended, resumed, and interleaved in a deterministic manner. This integration aims to adapt the coroutine scheduler—originally designed for graphics rendering tasks—to meet the real-time and safety-critical demands of autonomous driving. By doing so, the system can address the limitations of Apollo's existing batch-based scheduling, facilitating more dynamic task management and offering significant improvements in responsiveness and predictability.

The implementation will start with a pilot integration into a single, self-contained module within Apollo. This initial step serves as a proof-of-concept, focusing on adapting the LuisaCompute-coroutine scheduling to manage GPU threads effectively for one critical subsystem. Once the viability and benefits of this approach are validated—through detailed performance benchmarks and safety evaluations—the solution will be iteratively expanded to encompass additional modules across the platform. This staged rollout not only minimizes disruption to Apollo's existing architecture but also provides a controlled environment to fine-tune the integration, setting the stage for broader adoption of coroutine-based scheduling across the entire system.

Abbreviations

List of Figures

List of Tables